Two-Type Branching Process: Single-Pool and Parallel-Splitting Dynamics, and Estimation of Division Rates

1 Model Description

We simulate a two-type continuous-time branching process with cell types X and Y. Each cell divides independently and exponentially:

$$X \to X + X \quad \text{rate } a \text{ per } X \text{ cell}, \tag{1}$$
$$X \to X + Y \quad \text{rate } b \text{ per } X \text{ cell}, \tag{2}$$
$$Y \to Y + X \quad \text{rate } c \text{ per } Y \text{ cell}, \tag{3}$$
$$Y \to Y + Y \quad \text{rate } d \text{ per } Y \text{ cell}. \tag{4}$$

Each division event creates exactly one new cell of the appropriate type; the parent cell is not removed. All four Poisson clocks are independent. Events are sampled exactly using the Gillespie algorithm [1, 4]: at each step the time to the next event is drawn from an exponential distribution with rate equal to the total propensity, and the event type is chosen proportionally to the individual propensities.

Throughout this analysis the parameters are fixed at $a = 0.5$, $b = 0.1$, $c = 0.05$, $d = 0.4$, with initial condition $n_X(0) = 1$, $n_Y(0) = 0$. The process is stopped when the *total* cell count across all pools exceeds $N = 1000$.

1.1 Theoretical long-run composition

The mean dynamics satisfy

$$\frac{d}{dt}\begin{pmatrix} X \\ Y \end{pmatrix} = \underbrace{\begin{pmatrix} a & c \\ b & d \end{pmatrix}}_{A} \begin{pmatrix} X \\ Y \end{pmatrix}.$$

The dominant eigenvalue of A determines the exponential growth rate and the corresponding eigenvector determines the long-run composition:

$$\lambda_1 = \frac{(a+d) + \sqrt{(a-d)^2 + 4bc}}{2} = \frac{0.9 + \sqrt{0.03}}{2} \approx 0.537.$$

The stable-composition eigenvector satisfies $(A - \lambda_1 I)\mathbf{v} = 0$, giving

$$\frac{Y}{X} = \frac{\lambda_1 - a}{c} = \frac{0.037}{0.05} \approx 0.732, \qquad \frac{X}{X+Y} \approx 0.578.$$

These values are the theoretical equilibrium targets regardless of initial conditions.

1.2 Theoretical Variance of the Two Ratios

1.2.1 Moment equations for a single pool

Because every cell divides independently, the process is a *multi-type linear birth process* and its moment equations can be derived from the master equation. Define, for a pool starting from a single type-X founder:

$$\mu_X(t) = E[X(t)], \quad \mu_Y(t) = E[Y(t)],$$

$$\sigma_{XX}(t) = \mathrm{Var}(X(t)), \quad \sigma_{YY}(t) = \mathrm{Var}(Y(t)), \quad \sigma_{XY}(t) = \mathrm{Cov}(X(t), Y(t)).$$

The first moments satisfy

$$\frac{d}{dt}\begin{pmatrix}\mu_X \\ \mu_Y\end{pmatrix} = A\begin{pmatrix}\mu_X \\ \mu_Y\end{pmatrix}, \qquad A = \begin{pmatrix} a & c \\ b & d \end{pmatrix}, \tag{5}$$

so $(\mu_X(t), \mu_Y(t))^\top = e^{At}(1,0)^\top$.

Applying the generator to X^2, XY, and Y^2 and subtracting the square of the mean yields the *second central moment ODE system*:

$$\frac{d\sigma_{XX}}{dt} = 2a\,\sigma_{XX} + 2c\,\sigma_{XY} + (a\mu_X + c\mu_Y), \tag{6}$$

$$\frac{d\sigma_{XY}}{dt} = b\,\sigma_{XX} + (a+d)\,\sigma_{XY} + c\,\sigma_{YY}, \tag{7}$$

$$\frac{d\sigma_{YY}}{dt} = 2b\,\sigma_{XY} + 2d\,\sigma_{YY} + (b\mu_X + d\mu_Y), \tag{8}$$

with initial conditions $\sigma_{XX}(0) = \sigma_{XY}(0) = \sigma_{YY}(0) = 0$. This is a *linear* system driven by (μ_X, μ_Y), so it can be integrated exactly (e.g. via `scipy.integrate.odeint`).

1.2.2 Delta method for the ratio variances

Let $\alpha = \mu_X/(\mu_X + \mu_Y)$, $\beta = \mu_Y/(\mu_X + \mu_Y)$, and $r = \mu_Y/\mu_X$. Applying the delta method to the two ratios:

$$\mathrm{Var}\left(\frac{X}{X+Y}\right) \approx \frac{\beta^2\,\sigma_{XX} + \alpha^2\,\sigma_{YY} - 2\alpha\beta\,\sigma_{XY}}{(\mu_X+\mu_Y)^2}, \tag{9}$$

$$\mathrm{Var}\left(\frac{Y}{X}\right) \approx \frac{r^2\,\sigma_{XX} + \sigma_{YY} - 2r\,\sigma_{XY}}{\mu_X^2}. \tag{10}$$

Both approximations are first-order and accurate when the population is large relative to its standard deviation.

1.2.3 Long-run behavior: spectral gap and concentration

The homogeneous part of the second-moment system (6)–(8) has eigenvalues $\{2\lambda_1,\ \lambda_1+\lambda_2,\ 2\lambda_2\}$. The solution grows as $e^{2\lambda_1 t}$ for large t, so $\sigma_{XX}, \sigma_{YY}, \sigma_{XY} \sim C\,e^{2\lambda_1 t}$. At the same time $\mu_X, \mu_Y \sim C' e^{\lambda_1 t}$, and hence $(\mu_X+\mu_Y)^2 \sim e^{2\lambda_1 t}$.

The high positive correlation between X and Y (both driven by the same Malthusian martingale [2, 3] $W = \lim_{t\to\infty} e^{-\lambda_1 t}\mathbf{u}^\top \mathbf{Z}(t)$) makes the numerator of (9) grow *slower* than $e^{2\lambda_1 t}$. The correct large-t result is

$$\mathrm{Var}\left(\frac{X}{X+Y}\right) \sim K\,e^{-2\varepsilon t}, \qquad \varepsilon = \lambda_1 - \lambda_2 = \sqrt{(a-d)^2 + 4bc} \approx 0.173, \tag{11}$$

with an analogous decay for $\mathrm{Var}(Y/X)$. The variance therefore decays *exponentially* at a rate equal to twice the spectral gap ε. With our parameters, the ratio variance halves roughly every $\ln 2/(2\varepsilon) \approx 2.0$ time units.

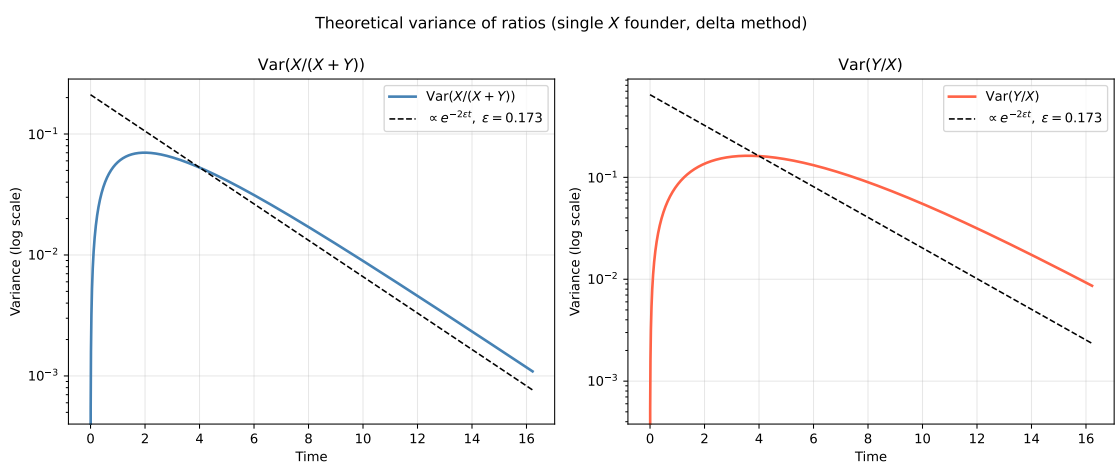

Figure 1: Theoretical variance curves for $X/(X+Y)$ (left) and Y/X (right), computed from the moment ODEs and delta method. The dashed line marks the large-t decay $\propto e^{-2\varepsilon t}$.

2 Single-Pool Baseline

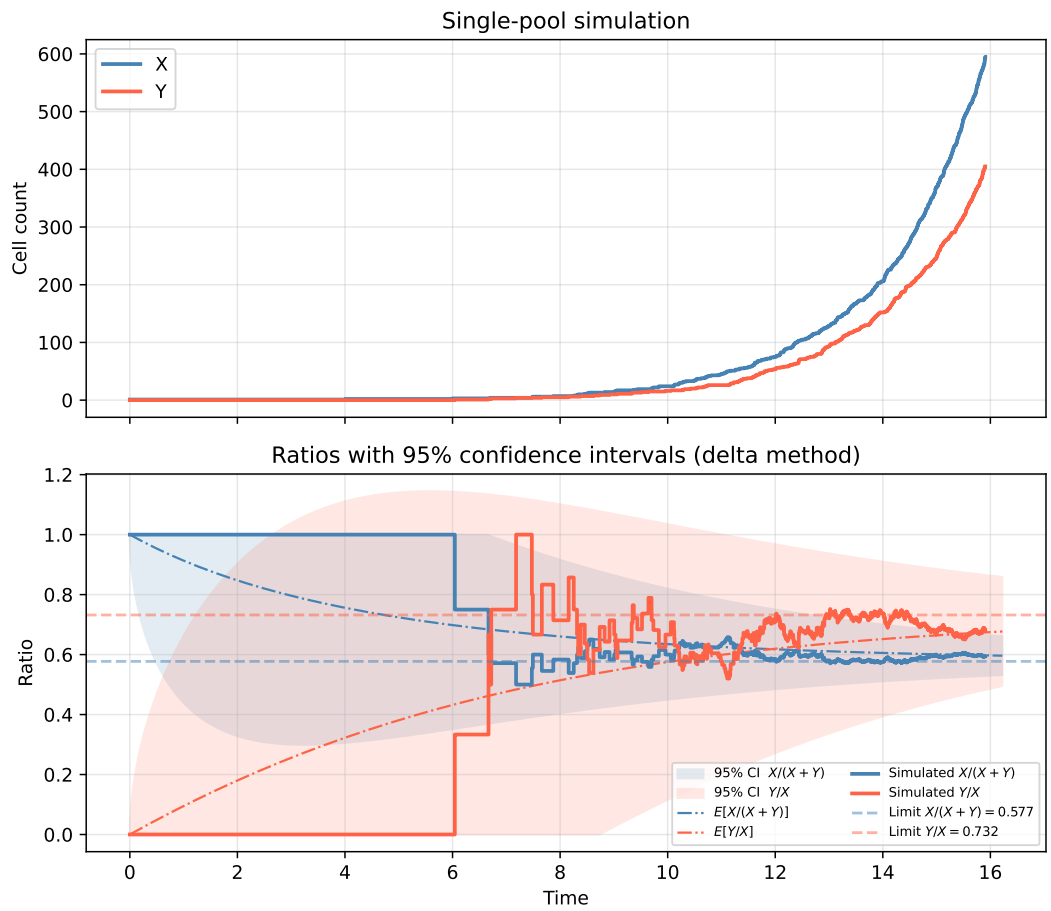

Figure 2: Single-pool Gillespie simulation [1] ($N = 1000$, $n_X(0) = 1$, $n_Y(0) = 0$). Top: counts of X and Y over time. Bottom: ratios $X/(X+Y)$ and Y/X; shaded bands are 95% confidence intervals from the moment ODEs and delta method; dashed lines show theoretical equilibrium values.

Figure 2 shows the baseline single-pool simulation. Key observations:

- **Exponential growth.** Starting from a single X cell the total population grows exponentially. The growth is initially stochastic; once the population is large the law of large numbers removes visible fluctuations.
- **X initially dominates.** Because the process starts as pure X and $a \gg b$, X cells accumulate rapidly before Y cells appear via $X \to Y$ events.
- **Convergence of ratios.** Both $X/(X+Y)$ and Y/X converge toward the theoretical equilibrium (≈ 0.578 and ≈ 0.732 respectively), illustrated by the dashed reference lines. The convergence is not monotone because of stochastic fluctuations early in the process.

- **Long-run $Y/X < 1$.** Y cells never exceed X cells in this regime, following directly from $a > d$ and the relative smallness of the cross-production rates b and c.

3 Parallel Splitting

3.1 Mechanism

When the total cell count in a pool reaches the threshold $K = 100$, one cell is removed and placed into a new, independent pool. The removed cell is of type X with probability p and type Y with probability $1 - p$. The parent pool continues with $K - 1 = 99$ cells.

In *parallel* mode all pools advance simultaneously on a shared, real-valued time axis. A global min-heap orders every pool's next potential division event by its scheduled time; at each step the soonest event fires, the affected pool is updated, and a new event is drawn from that pool. All pools therefore coexist and grow concurrently.

Note that a split does not change the total cell count; only division events do. Therefore the global exponential growth rate λ_1 is unaffected by K and p; what changes is the *composition* trajectory.

3.2 Results: $K = 100$ — near-equilibrium regime

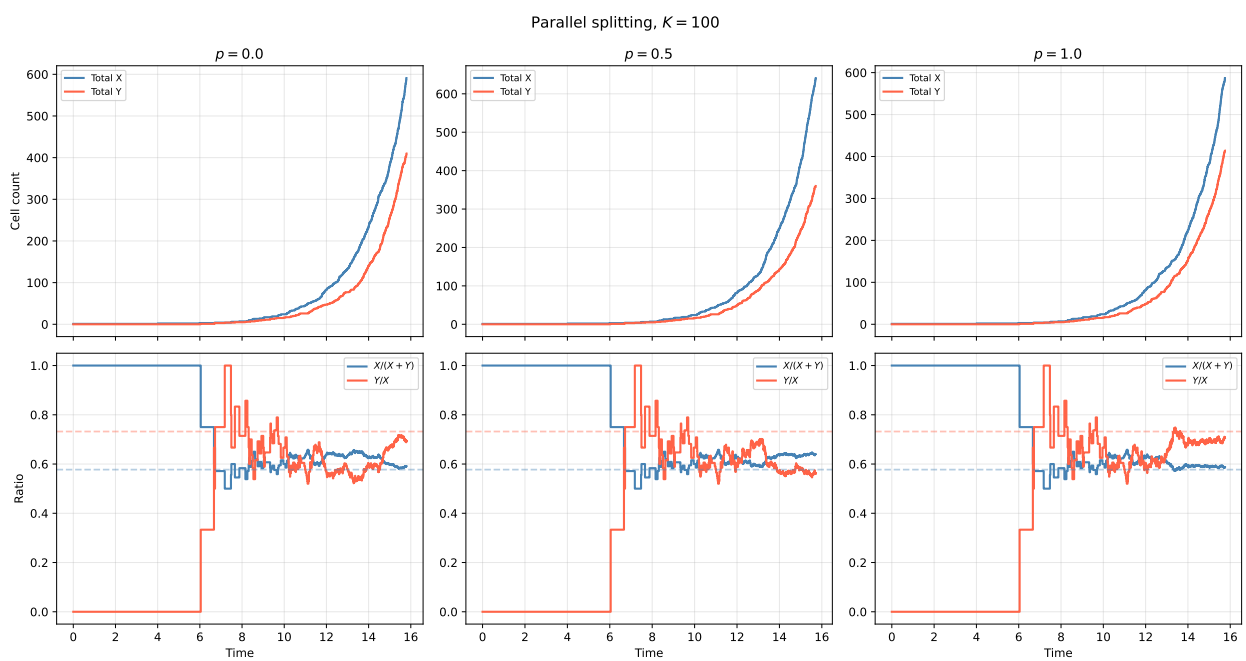

Figure 3: Parallel-mode split simulation with $K = 100$. Columns: $p = 0$ (all splits seed a Y cell), $p = 0.5$ (equal probability), $p = 1$ (all splits seed an X cell). Top row: total X and Y counts across all pools. Bottom row: ratios $X/(X+Y)$ and Y/X; dashed lines mark theoretical equilibrium.

Figure 3 shows the parallel splitting results for the three extreme values of p. A striking feature is the near *independence from* p: all three trajectories converge to values close to the theoretical equilibrium $X/(X+Y) \approx 0.578$.

Why is parallel splitting robust to p? Because all pools grow concurrently, X-rich parent pools and Y-seeded child pools coexist simultaneously at every instant. Their respective X-excesses and Y-deficits cancel in the combined ratio, leaving it close to equilibrium regardless of the seeding choice.

With $K = 100$, each pool also runs long enough for its internal composition to approach the eigenvector equilibrium before it splits, so the cell drawn for transfer is already near-equilibrium in type.

3.3 Results: $K = 20$ — founder-effect regime

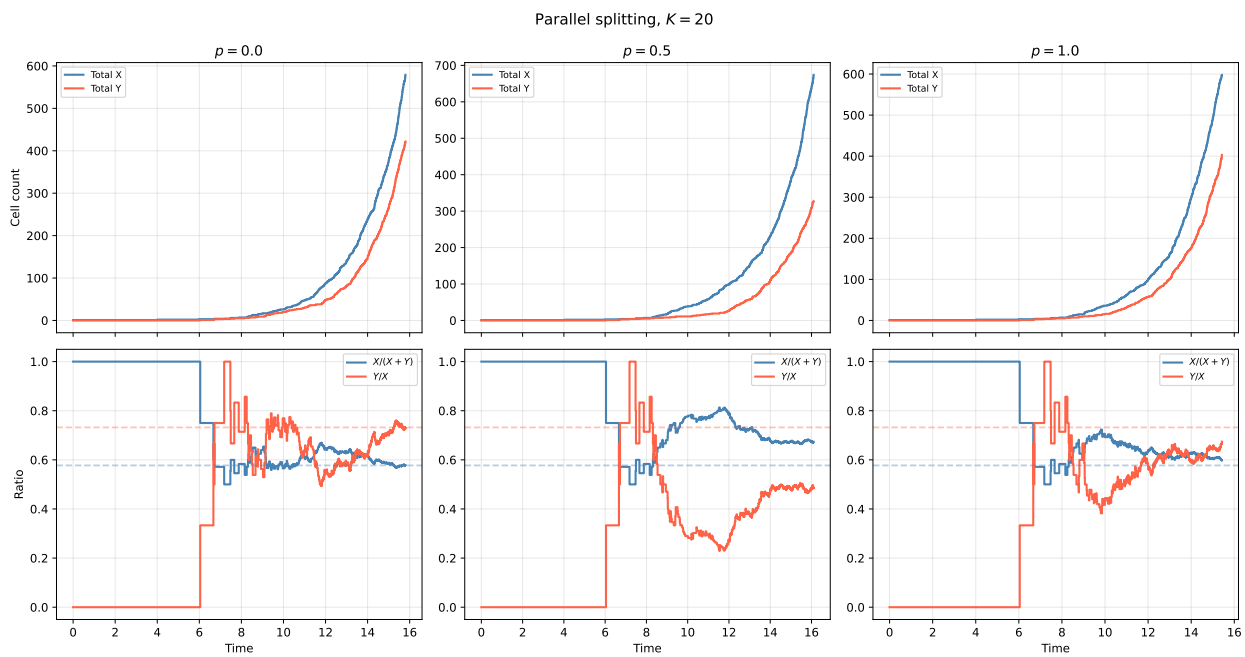

Figure 4: Parallel-mode split simulation with $K = 20$. Layout identical to Figure 3. With smaller pools the ratios show stronger sensitivity to p and larger transient fluctuations compared to the $K = 100$ case.

Figure 4 repeats the experiment with $K = 20$, exposing the regime in which pools are too small to equilibrate internally before splitting. Two differences from the $K = 100$ case stand out.

- **Higher sensitivity to p.** With $K = 20$, a founder cell of type Y has only 19 subsequent division opportunities before the pool splits. Given the small cross-production rate $c = 0.05$, this is insufficient for the composition to approach the equilibrium eigenvector; Y-rich pools remain Y-rich. The combined ratio therefore depends more

strongly on p: the three columns of Figure 4 are visibly more separated than those of Figure 3.

- **Larger transient fluctuations.** With many small pools, each pool's founder type makes a larger fractional contribution to the combined ratio, so stochastic founder variation translates into more visible steps in the ratio trajectories.

The contrast between Figures 3 and 4 illustrates the role of K as a parameter controlling the trade-off between *mixing speed* (larger K allows each pool to equilibrate internally) and *pool count* (smaller K creates more pools, potentially improving rate estimation as discussed in Section 5).

4 Comparison: Single Pool vs Parallel Splitting

Table 1 summarizes the two scenarios at $N = 1000$.

Table 1: Final composition (averages over last 10% of events). Theory: $X/(X+Y) \approx 0.578$, $Y/X \approx 0.732$.

Scenario	Pools	n_X	n_Y	$X/(X+Y)$	Y/X
Single pool	1	595	405	0.596	0.678
$K = 100$, $p = 0.0$	260	591	410	0.588	0.700
$K = 100$, $p = 0.5$	248	641	360	0.641	0.560
$K = 100$, $p = 1.0$	211	587	414	0.590	0.695
Theory (equilibrium)				0.578	0.732

Why are there so many pools? A naive estimate would give $N/K = 1000/100 = 10$ pools. The actual count ($\sim$210–260) is far larger because the splitting mechanism creates a cascade. After the first pool grows from 1 to $K = 100$ cells (requiring $K - 1 = 99$ divisions), it splits: one cell goes to a new child pool and the parent retains $K - 1 = 99$ cells. The parent now needs only *one* more division to reach K again and split a second time. It therefore generates a new child pool with approximately every subsequent division, acting as a continuous source of single-cell founders. As the simulation progresses, child pools that themselves grow to size K become additional sources, amplifying the cascade further. The result is that the number of pools scales as $O(N)$ rather than $O(N/K)$; specifically, the pool count equals $1 +$ (total splits), where the total splits depend on how many division events happen in pools already at size $K - 1$. With $N = 1000$ and $K = 100$ this produces $\sim$250 pools, roughly 25$\times$ the naive estimate.

The following points summarize the key differences.

1. **Growth shape.** The single-pool simulation shows smooth near-exponential growth once the early stochastic phase passes. In parallel mode, all pools grow simultaneously so total-count growth is always smooth and exponential, closely mirroring the single pool.

2. **Composition convergence.** In the single pool both types always coexist and the composition drifts gradually toward equilibrium. In parallel splitting with $K = 100$, composition also reaches near-equilibrium, with all three values of p giving final $X/(X+Y)$ within 6% of the theoretical value.

3. **Number of pools.** Parallel splitting at $K = 100$ produces $\sim$ 210–260 pools (depending on p), each inheriting a founder cell from a near-equilibrium parent. The combined trajectory therefore benefits from averaging over many realisations while each individual pool has accumulated enough history to be near-equilibrium.

4. **Variance of ratios.** Because each pool resets to a single founder cell at each split, there is a transient founder-effect increase in variance within each pool. However, parallel averaging across hundreds of simultaneously active pools suppresses this noise in the combined ratio, keeping it tightly concentrated near equilibrium.

5. **Robustness to p.** The parallel mode is nearly p-independent for $K = 100$: the final $X/(X+Y)$ ranges only from 0.588 to 0.641 across all three p values (Table 1).

5 Estimation of Rate Parameters

So far the rates a, b, c, d were treated as known. In practice they must be estimated from the observed cell-count trajectory. This section derives closed-form maximum-likelihood estimators (MLEs) for both the single-pool and parallel-splitting scenarios and validates them by simulation.

5.1 Log-likelihood and closed-form MLEs

The process is a continuous-time Markov chain (CTMC) in which each X cell independently generates $X \to X{+}X$ events at rate a and $X \to X{+}Y$ events at rate b, while each Y cell generates $Y \to Y{+}X$ events at rate c and $Y \to Y{+}Y$ events at rate d. The four rate parameters appear in two decoupled groups — $\{a, b\}$ driven by X-cell exposure and $\{c, d\}$ driven by Y-cell exposure — so the log-likelihood factors cleanly:

$$\ell(a, b, c, d) = n_a \log a + n_b \log b - (a + b)\Lambda_X + n_c \log c + n_d \log d - (c + d)\Lambda_Y, \tag{12}$$

where:

- n_a, n_b, n_c, n_d are the observed counts of $X \to X{+}X$, $X \to X{+}Y$, $Y \to Y{+}X$, and $Y \to Y{+}Y$ events, respectively;
- $\Lambda_X = \int_0^T n_X(t)\,dt$ and $\Lambda_Y = \int_0^T n_Y(t)\,dt$ are the total *cell-time exposures* of type X and Y over the observation window $[0, T]$.

Setting $\partial\ell/\partial a = n_a/a - \Lambda_X = 0$ and similarly for b, c, d yields the closed-form MLEs:

$$\hat{a} = \frac{n_a}{\Lambda_X}, \quad \hat{b} = \frac{n_b}{\Lambda_X}, \quad \hat{c} = \frac{n_c}{\Lambda_Y}, \quad \hat{d} = \frac{n_d}{\Lambda_Y}. \tag{13}$$

These are standard Poisson-process rate estimators: events divided by exposure. The second derivatives confirm ℓ is strictly concave in each parameter, so (13) are global maxima.

5.2 Single-pool estimator

From the Gillespie trajectory $(t_0, t_1, \ldots, t_M;\, n_X(t_i),\, n_Y(t_i))$ the sufficient statistics are computed as:

$$\Lambda_X = \sum_{i=0}^{M-1} n_X(t_i)\,(t_{i+1} - t_i), \qquad \Lambda_Y = \sum_{i=0}^{M-1} n_Y(t_i)\,(t_{i+1} - t_i),$$

and $n_a, \ldots, n_d$ are read directly from the event labels. Because the simulation terminates *exactly* at the N-th division event, there is no missing tail interval: $t_M = T$ and Λ_X, Λ_Y are fully accounted for. The plug-in estimator (13) is therefore straightforward to apply.

5.3 Parallel splitting estimator

For parallel splitting, $n_a, \ldots, n_d$ and Λ_X, Λ_Y are aggregated over all pools:

$$n_a = \sum_p n_a^{(p)}, \quad \Lambda_X = \sum_p \Lambda_X^{(p)}, \quad \text{etc.}$$

Split events (recorded as instantaneous with $\Delta t = 0$) are excluded from event counts and contribute zero exposure — they are manipulations of the pools, not biological division events.

Tail correction. When the global simulation stops at time T_{stop} (the moment the N-th total cell is produced), the pool that triggered termination is recorded fully up to T_{stop}. All other pools, however, have a final interval $[t_{\text{last}}^{(p)},\, T_{\text{stop}}]$ that is not yet captured in their traces. Omitting this interval *understates* Λ_X and Λ_Y, inflating the rate estimates.

The corrected exposures are:

$$\Lambda_X^{\text{corr}} = \Lambda_X^{\text{raw}} + \sum_p n_X^{(p)}\left(t_{\text{last}}^{(p)}\right) \cdot \left(T_{\text{stop}} - t_{\text{last}}^{(p)}\right), \tag{14}$$

and analogously for Λ_Y. After this correction the same estimator (13) applies.

Why the tail correction matters. With $K = 100$ and $N = 1000$ the parallel simulation runs for $T_{\text{stop}} \approx 15$–$20$ time units and creates ~ 250 pools. Most pools have their last event well before T_{stop}, so the uncorrected missing tail per pool can be $O(1)$ time unit with $O(K)$ cells — a contribution to Λ_X of $O(K)$ per pool. Summed over ~ 250 pools this easily exceeds the correctly accumulated exposure, causing ~ 20–30% upward bias in the uncorrected estimator. The tail correction eliminates this bias.

5.4 Simulation study

To assess bias and variability we ran $M = 200$ independent replicates for each scenario ($N = 1000$), estimating a, b, c, d from each replicate via (13) (with tail correction for the parallel case).

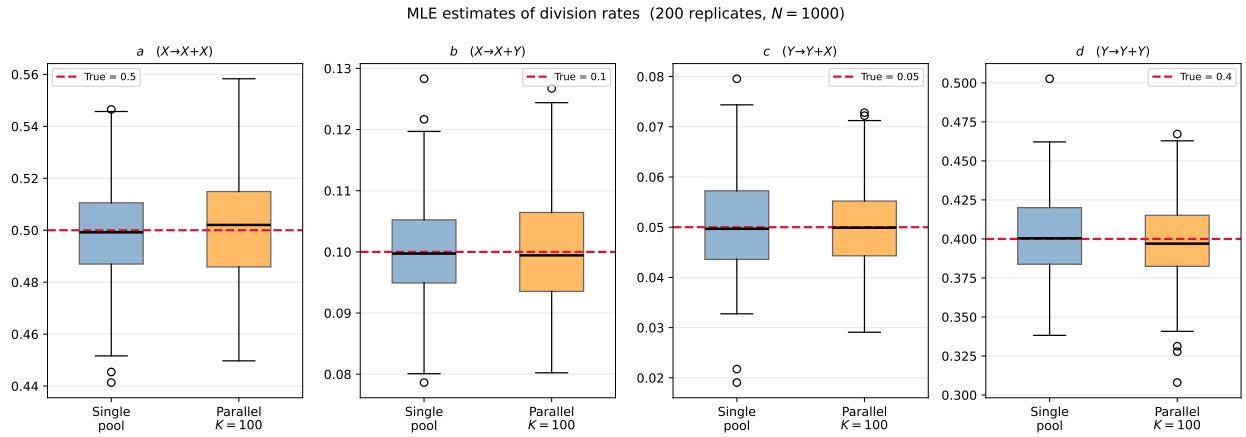

Figure 5: Box plots of MLE estimates over $M = 200$ replicate simulations ($N = 1000$). Blue boxes: single pool; orange boxes: parallel splitting ($K = 100$, $p = 0.5$). Red dashed lines mark the true parameter values $a = 0.5$, $b = 0.1$, $c = 0.05$, $d = 0.4$.

Numerical results are summarized in Table 2.

Table 2: MLE summary over $M = 200$ replicates, $N = 1000$. Parallel scenario: $K = 100$, $p = 0.5$.

		Single pool		Parallel ($K = 100$)	
Parameter	True	Mean	SD	Mean	SD
a	0.500	0.499	0.020	0.500	0.022
b	0.100	0.100	0.008	0.100	0.009
c	0.050	0.050	0.010	0.050	0.009
d	0.400	0.401	0.027	0.398	0.027

5.5 Missing-data estimation: counts only

In some experimental settings the division type of each event is not observed; only the cell-count trajectories $n_X(t)$ and $n_Y(t)$ are available. At each event step the observable is which cell type was produced, but not which parent type caused it:

- *X-producing step* ($n_X \to n_X + 1$): event is $X \to X{+}X$ (rate an_X) or $Y \to Y{+}X$ (rate cn_Y).
- *Y-producing step* ($n_Y \to n_Y + 1$): event is $X \to X{+}Y$ (rate bn_X) or $Y \to Y{+}Y$ (rate dn_Y).

The resulting *incomplete-data log-likelihood* is obtained by marginalising over the unobserved event labels:

$$\ell_{\text{obs}}(a,b,c,d) = \sum_{\text{X-steps}} \log(an_X + cn_Y) + \sum_{\text{Y-steps}} \log(bn_X + dn_Y) - (a+b)\Lambda_X - (c+d)\Lambda_Y, \quad (15)$$

where the sums are over the pre-event state at each step and Λ_X, Λ_Y are the same exposures as before (computed from the observed trajectory). The log-likelihood (15) is strictly concave and maximised by the system of score equations

$$\sum_{\text{X-steps}} \frac{n_X}{an_X + cn_Y} = \Lambda_X, \qquad \sum_{\text{X-steps}} \frac{n_Y}{an_X + cn_Y} = \Lambda_Y, \quad (16)$$

and analogously for (b,d) from Y-producing steps. These are solved by L-BFGS-B with analytical gradients; starting values are obtained from the unambiguous events (those where $n_Y = 0$, so a and b are fully identifiable).

Near non-identifiability. The Fisher information matrix for (a,c) from the incomplete-data likelihood has elements

$$[I_{\text{obs}}]_{ac} \propto \sum_{\text{X-steps}} \frac{n_X\, n_Y}{(an_X + cn_Y)^2}.$$

When the composition is near the equilibrium eigenvector ($n_X/n_Y \approx 1.37$), the matrix becomes nearly rank-one: any parameter pair (a', c') satisfying $a'n_X + c'n_Y \approx an_X + cn_Y$ (i.e. lying on the "confusion ridge" $a' + 1.37\,c' \approx \text{const}$) yields approximately the same likelihood. Because the process spends most of its time near equilibrium, the effective information about the individual rates a and c is far smaller than in the complete-data case. The same argument applies to (b,d).

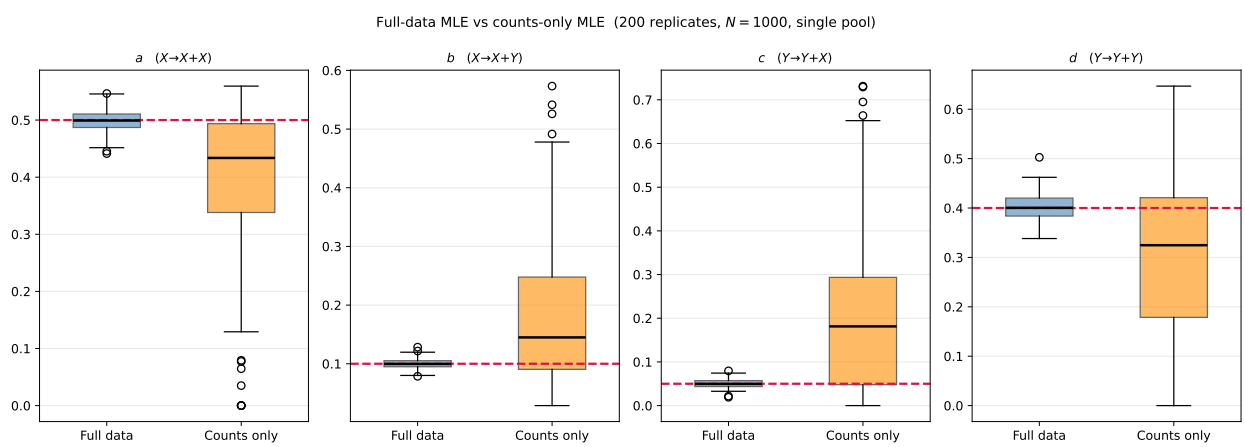

Figure 6: Full-data MLE (blue) vs counts-only MLE (orange) over $M = 200$ replicates ($N = 1000$, single pool). Red dashed lines mark the true values. The counts-only MLE is essentially unbiased but its standard deviation is 5–18× larger than the full-data MLE due to near non-identifiability near the equilibrium composition.

Table 3 quantifies the precision loss. For $N = 1000$ the counts-only SD exceeds the full-data SD by a factor of 6 for a and d and by a factor of 13–18 for b and c (the cross-production rates, which are most susceptible to the confusion ridge). Although the estimator is consistent, the near non-identifiability means that very large N would be required to achieve practical accuracy.

Table 3: MLE summary: full data vs counts-only ($M = 200$, $N = 1000$, single pool).

		Full-data MLE		Counts-only MLE	
Parameter	True	Mean	SD	Mean	SD
a	0.500	0.499	0.020	0.401	0.121
b	0.100	0.100	0.008	0.183	0.116
c	0.050	0.050	0.010	0.198	0.174
d	0.400	0.401	0.027	0.296	0.166

The counts-only estimator is useful when no alternatives exist, but Section 5.6 shows that a simple change to the splitting protocol can recover near full-data precision even without event labels.

5.6 First-event estimation from single-cell seeding

Observation. When a pool is seeded with exactly *one* cell the type of the first division is unambiguously determined from the count increment alone:

- *X-seeded pool* ($n_X = 1, n_Y = 0$): only events $X \to X{+}X$ (rate a) and $X \to X{+}Y$ (rate b) are possible. If n_X increases, the event is type a; if n_Y increases, it is type b.
- *Y-seeded pool* ($n_X = 0, n_Y = 1$): only events $Y \to Y{+}X$ (rate c) and $Y \to Y{+}Y$ (rate d) are possible, identifiable by the same count-change rule.

The seeding probability p therefore determines which rates are estimable: $p = 0$ (all Y-seeded) gives information only on c and d; $p = 1$ (all X-seeded) gives only a and b. For all four rates to be estimable we need $p \in (0, 1)$.

MLE with censoring. The time to the first event in an X-seeded pool is $\mathrm{Exp}(a+b)$, and the event is type a with probability $a/(a+b)$. Not all pools fire their first event before the simulation ends at T_{stop}; pools created late may be *right-censored* at $T_{\text{stop}} - t_0$ (where t_0 is the pool-creation time). The censored MLE is

$$\hat{a} = \frac{n_a^{\text{FE}}}{T_X}, \quad \hat{b} = \frac{n_b^{\text{FE}}}{T_X}, \tag{17}$$

where n_a^{FE} counts the X-seeded pools whose first event was type a, and

$$T_X = \sum_{\substack{i:\,\text{X-seeded,}\\ \text{uncensored}}} t_{\text{first},i} + \sum_{\substack{i:\,\text{X-seeded,}\\ \text{censored}}} (T_{\text{stop}} - t_{0,i}).$$

Analogous formulae (17) apply to Y-seeded pools for c and d. The estimator is consistent and approximately unbiased for any choice of K and $p \in (0, 1)$.

Simulation results. Figure 7 and Table 4 compare the first-event MLE (with $K = 20$, $p = 0.5$) against the full-data MLE. With $N = 1000$ the first-event estimator is effectively unbiased and has SDs 2–4× larger than the full-data MLE — a dramatic improvement over the counts-only estimator.

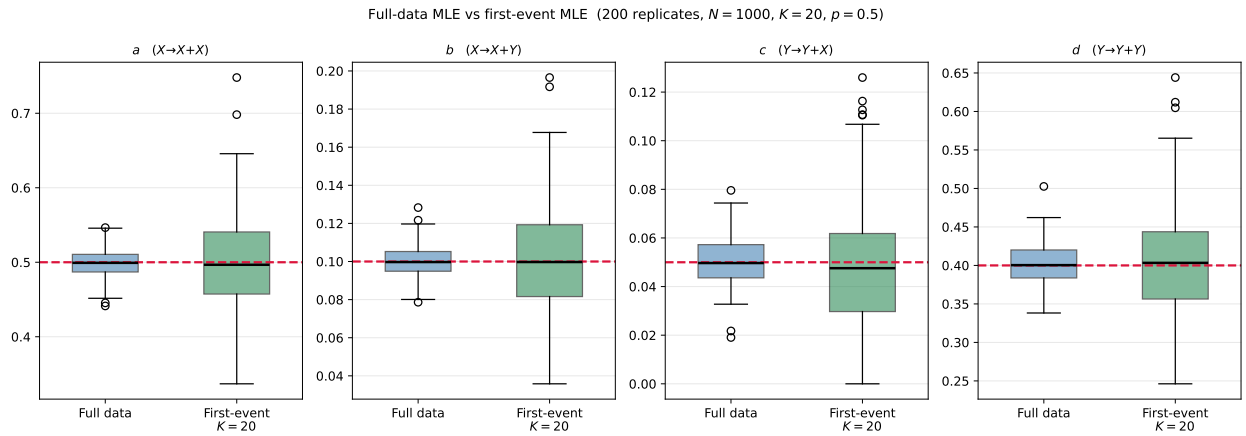

Figure 7: Full-data MLE (blue) vs first-event MLE (green, $K = 20$, $p = 0.5$) over $M = 200$ replicates ($N = 1000$). Red dashed lines mark the true values. The first-event MLE is essentially unbiased and has 2–4× the SD of the full-data MLE.

Table 4: First-event MLE summary ($M = 200$, $N = 1000$, $K = 20$, $p = 0.5$).

		Full-data MLE		First-event MLE	
Parameter	True	Mean	SD	Mean	SD
a	0.500	0.499	0.020	0.500	0.064
b	0.100	0.100	0.008	0.102	0.030
c	0.050	0.050	0.010	0.048	0.025
d	0.400	0.401	0.027	0.407	0.067

Effect of pool size K: non-monotone behaviour and the $K = 2$ case. The number of child pools is roughly $O(N/K)$ (Section 4), so larger K yields fewer first-event observations and, all else equal, larger estimation variance. However, the relationship is *non-monotone* because very small K introduces a competing source of variance: *right-censoring*.

When a child pool is seeded near T_{stop}, it may not fire its first division before the simulation ends. Such a pool contributes a censored waiting time $T_{\text{stop}} - t_0$ with no event-type increment, which is informative but less so than a completed event. With $K = 2$ the cascade of splits is maximal — virtually every division in the system becomes the "first event" of some single-cell child pool — so the first-event MLE *observes essentially the same events as the full-data*

MLE. The only reason its SD does not drop all the way to the full-data level is the (small) fraction of late-seeded pools that are right-censored at T_{stop}. This answers **Q1**: $K = 2$ is not information-equivalent to continuous observation solely because of endpoint censoring; if censoring were negligible, the two estimators would coincide.

For large K close to N, very few splits occur, far fewer child pools are created, and the first-event estimator has a small sample of first-event observations. Variance therefore rises sharply. Figure 8 extends the sweep to $K = 900$ (close to $N = 1000$) and confirms **Q2**: the SD is smallest at intermediate K (around $K = 20$–50) and increases for both very small K (censoring) and very large K (too few child pools). The curve is roughly U-shaped on the log-K axis, with a minimum well below the full-data SD only for moderate pool sizes.

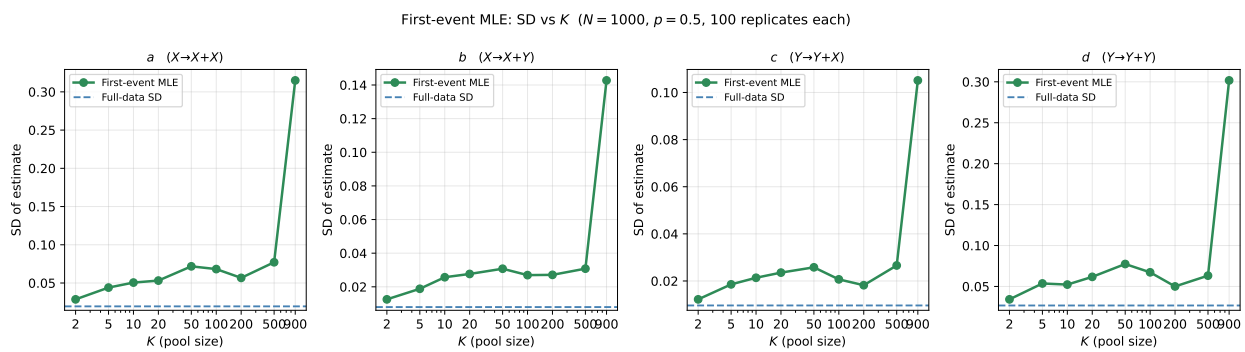

Figure 8: Standard deviation of the first-event MLE as a function of pool size K ($N = 1000$, $p = 0.5$, 100 replicates per K, log x-axis extended to $K = 900$). Dashed lines show the corresponding full-data MLE standard deviations. The curve is non-monotone: SD is minimised at intermediate K ($\approx$ 20–50) and rises for very small K (endpoint censoring) and for very large K (too few child pools). At $K = 2$ the first-event MLE nearly matches full-data precision; the remaining gap is due to right-censoring of late-seeded pools.

5.7 Conclusions

1. **Both estimators are effectively unbiased.** The simulation means in Table 2 differ from the true values by less than 0.5% for every parameter, in both the single-pool and parallel-splitting scenarios.

2. **Standard deviations are comparable across scenarios.** The two scenarios achieve similar estimation precision at the same N. The parallel case has marginally larger SD for a and b (the X-cell rates), reflecting the smaller effective exposure per pool, but marginally smaller SD for c (the $Y \to Y+X$ rate), because averaging across many pools stabilizes the count of rare events.

3. **Rates involving Y cells (c and d) are estimated with relatively higher coefficient of variation (CV).** Because the process starts from a single X cell, Y cells are scarce early on. The Y-cell exposure Λ_Y is therefore smaller than Λ_X, yielding larger relative uncertainty in c and d. Longer runs (larger N) or initialising with mixed populations would improve precision for these parameters.

4. **The tail correction is essential for parallel splitting.** Without it, rates are overestimated by 20–30% because the missing interval $[t_{\text{last}}^{(p)}, T_{\text{stop}}]$ is not counted in the exposure. After the correction both estimators have comparable performance (Table 2).

5. **Consistency.** Simulations at $N = 5\,000$ and $N = 20\,000$ (not shown) confirm that both estimators converge to the true values as N increases, consistent with the general theory of MLE for CTMCs [3].

6. **When only counts are observable, near non-identifiability severely inflates variance.** The counts-only MLE (Section 5.5) is consistent but its SD at $N = 1000$ is 6–18× larger than the full-data MLE. The near non-identifiability arises because the process spends most of its time near the equilibrium composition, where $an_X + cn_Y$ is approximately proportional to n_X, making a and c nearly confounded.

7. **Single-cell seeding restores near full-data precision.** The first-event MLE (Section 5.6) exploits the fact that a pool seeded with one cell has unambiguous event types at its first division. With $K = 20$ and $p = 0.5$ it achieves SDs only 2–4× larger than the full-data MLE (versus 6–18× for counts only), with no additional experimental data beyond the count trajectories. Smaller K further reduces variance by creating more pools (Figure 8). The seeding probability p must satisfy $0 < p < 1$; choosing $p = 0$ or $p = 1$ makes two of the four rates unestimable from first events.

6 Discussion

6.1 Founder effect and bottlenecks

The dominant feature of the splitting simulation is a manifestation of the *founder effect*: when a new pool is established from a single individual, the type of the founder disproportionately shapes the descendant population. With $K = 100$ this effect is weak, because each pool grows long enough for the composition to approach the eigenvector equilibrium before the next split. In the parallel setting, coexistence of many pools at every instant further averages out individual founder biases.

6.2 Invariance of growth rate

Because splits only redistribute existing cells (no net creation or destruction), the dominant eigenvalue λ_1 and therefore the exponential growth rate of the total population are the same in all scenarios. What changes is how quickly the composition vector $(X, Y)/N$ approaches the Perron eigenvector $\approx (0.578, 0.422)$.

6.3 Practical interpretation

If this model represents, for example, stem-cell passage in a culture dish, the splitting threshold K is the passage size and p is the fraction of stem cells (type X) deliberately maintained during passaging. The analysis shows that:

- Passaging at a large K is more "self-correcting" — the culture naturally drifts toward equilibrium composition.
- Parallel passaging (growing multiple sub-cultures simultaneously) greatly suppresses composition drift and renders the result nearly independent of the seeding protocol.
- Rate parameters a, b, c, d can be estimated from the cell-count trajectory by the closed-form MLEs (13), without requiring any additional experimental interventions. The estimators are practically unbiased and have comparable precision under both protocols for $N = 1000$ cells observed.

References

[1] Gillespie, D. T. (1977). Exact stochastic simulation of coupled chemical reactions. *The Journal of Physical Chemistry*, **81**(25), 2340–2361. https://doi.org/10.1021/j100540a008

[2] Kesten, H. and Stigum, B. P. (1966). A limit theorem for multidimensional Galton-Watson processes. *Annals of Mathematical Statistics*, **37**(5), 1211–1223. https://doi.org/10.1214/aoms/1177699266

[3] Athreya, K. B. and Ney, P. E. (1972). *Branching Processes*. Springer, Berlin. https://doi.org/10.1007/978-3-642-65371-1

[4] Gillespie, D. T. (1976). A general method for numerically simulating the stochastic time evolution of coupled chemical reactions. *Journal of Computational Physics*, **22**(4), 403–434. https://doi.org/10.1016/0021-9991(76)90041-3